

CSIT115 Data Management and Security

# Database Design Quality

Dr Janusz R. Getta

School of Computing and Information Technology -  
University of Wollongong

# Database Design Quality

## Outline

Why not ONE BIG TABLE !?

Where is a problem ?

Insertion test

Join test

Deletion test


Update test

Good design guidelines

# Why not ONE BIG TABLE !?

Let us consider the following database domain:

- A **customer** is described by a unique **customer number**, **first**, and **last name**
- **Customers** submit **orders**. An **order** is described by a **unique order number** and **order date**
- **Orders** consist of **lines**. A **line** contains information about a **name of ordered item**, **price per single item**, and **total number of ordered items**



# Why not ONE BIG TABLE !?

Logical design provides the following relational schemas:

Relational schemas

```
CUSTOMER(cnumber, first-name, last-name)
PRIMARY KEY = (cnumber)

ORDERS(onumber, odate, cnumber) PRIMARY KEY = (onumber)
FOREIGN KEY = (cnumber) REFERENCES CUSTOMER(cnumber)

LINE(onumber, lnumber, item, price total)
PRIMARY KEY = (onumber, lnumber)
FOREIGN KEY = (onumber) REFERENCES ORDERS(onumber)
```

Why not one relational schema ?

Big relational schema

```
CUSTOMER(cnumber, first-name, last-name, onumber, odate, cnumber,
          onumber, lnumber, item, price total)
PRIMARY KEY = (cnumber, onumber, lnumber)
```

# Why not ONE BIG TABLE !?

Insertion of information about one customer who submitted 2 orders such that each order consists several lines reveals a problem !

Big relational table

cnumber	fname	lname	onumber	odate	lnumber	item	price	total
7	James	Bond	7	2017-01-01	1	bolt	23.04	5
7	James	Bond	7	2017-01-01	2	screw	29.01	3
7	James	Bond	7	2017-01-01	3	nut	4.55	2
7	James	Bond	8	2018-01-01	1	bolt	23.04	1
7	James	Bond	8	2018-01-01	2	screw	23.04	1
7	James	Bond	8	2018-01-01	3	nut	23.04	2
7	James	Bond	8	2018-01-01	4	lock	23.04	1

A **number**, **first name**, and **last name** of a customer is repeated as many times as the total number of different items purchased in all orders and **order number** is repeated together with **order date** as many times as the total number of different items purchased in an order

# Why not ONE BIG TABLE !?

A multitable design does not have such a problem:

CUSTOMER(cnumber, first-name, last-name)

PRIMARY KEY = (cnumber)

CUSTOMER schema

cnumber	fname	lname
7	James	Bond

CUSTOMER table

ORDERS(onumber, odate, cnumber) PRIMARY KEY = (onumber)

FOREIGN KEY = (cnumber) REFERENCES CUSTOMER(cnumber)

ORDERS schema

onumber	odate	cnumber
7	2017-01-01	7
8	2018-01-01	7

ORDERS table

# Why not ONE BIG TABLE !?

A multitable design does not have such a problem:

```
LINE(onenumber, lnumber, item, price total)
PRIMARY KEY = (onenumber, lnumber)
FOREIGN KEY = (onenumber) REFERENCES ORDERS(onenumber)
```

LINE schema

```
+-----+-----+-----+-----+-----+
| onumber | lnumber | item  | price | total |
+-----+-----+-----+-----+-----+
|      7 |      1 | bolt  | 23.04 |    5 |
|      7 |      2 | screw | 29.01 |    3 |
|      7 |      3 | nut   |  4.55 |    2 |
|      8 |      1 | bolt  | 23.04 |    1 |
|      8 |      2 | screw | 23.04 |    1 |
|      8 |      3 | nut   | 23.04 |    2 |
|      8 |      4 | lock  | 23.04 |    1 |
+-----+-----+-----+-----+-----+
```

LINE table

# Database Design Quality

## Outline

Why not ONE BIG TABLE !?

Where is a problem ?

Insertion test

Join test

Deletion test






















Update test

Good design guidelines



# Where is a problem ?

Why do we get redundancies in an incorrectly designed relational table ?

TABLE_NAME			
COLUMN_1	COLUMN_2	...	COLUMN_N
		...	
		...	
		...	
		...	
		...	
		...	
		...	

Data dependencies:

- If COLUMN\_1 is green then COLUMN\_2 is red
- If COLUMN\_1 is blue then COLUMN\_2 is yellow
- If COLUMN\_1 is orange then COLUMN\_2 is red
- For any colour x if COLUMN\_1 is x then COLUMN\_2 is y

# Where is a problem ?

Data dependencies can be represented as a separate relational table ...

TABLE\_1

COLUMN_1	COLUMN_2

... and `COLUMN_2` can be removed from the original table

TABLE\_2

COLUMN_1	...	COLUMN_N
	...	
	...	
	...	
	...	
	...	
	...	
	...	

# Where is a problem ?

Do data dependencies exist in BIG TABLE ?

Big relational table

cnumber	fname	lname	onumber	odate	lnumber	item	price	total
7	James	Bond	7	2017-01-01	1	bolt	23.04	5
7	James	Bond	7	2017-01-01	2	screw	29.01	3
7	James	Bond	7	2017-01-01	3	nut	4.55	2
7	James	Bond	8	2018-01-01	1	bolt	23.04	1
7	James	Bond	8	2018-01-01	2	screw	23.04	1
7	James	Bond	8	2018-01-01	3	nut	23.04	2
7	James	Bond	8	2018-01-01	4	lock	23.04	1

Data dependencies:

- If `cnumber` = 7 then `fname` = James
- If `cnumber` = 7 then `lname` = Bond
- For any customer number x if `cnumber` = x then `fname` = y and `lname` = z

# Where is a problem ?

Do data dependencies exist in BIG TABLE ?

Big relational table

cnumber	fname	lname	onumber	odate	lnumber	item	price	total
7	James	Bond	7	2017-01-01	1	bolt	23.04	5
7	James	Bond	7	2017-01-01	2	screw	29.01	3
7	James	Bond	7	2017-01-01	3	nut	4.55	2
7	James	Bond	8	2018-01-01	1	bolt	23.04	1
7	James	Bond	8	2018-01-01	2	screw	23.04	1
7	James	Bond	8	2018-01-01	3	nut	23.04	2
7	James	Bond	8	2018-01-01	4	lock	23.04	1

More data dependencies:

- If `onumber` = 7 then `odate` = 2017-01-01
- If `onumber` = 8 then `odate` = 2018-01-01
- For any order number x if `onumber` = x then `odate` = y

# Database Design Quality

## Outline

Why not ONE BIG TABLE !?

Where is a problem ?

Insertion test

Join test

Deletion test

Update test

Good design guidelines

# Insertion test

How to verify if a relational schema is designed in a correct way ?

We try to insert few rows such that it is possible to create redundancies

For example, we insert few rows into a relational table **ROOM**

```
ROOM(bldgnum, name, roomnum, area)
```

```
PRIMARY KEY = (bldgnum, roomnum)
```

ROOM schema

```
+-----+-----+-----+-----+
| bldgnum | bldgname | roomnum | area |
+-----+-----+-----+-----+
|      3 | SCIT    |     210 |   20 |
+-----+-----+-----+-----+
```

ROOM table

```
+-----+-----+-----+-----+
| bldgnum | bldgname | roomnum | area |
+-----+-----+-----+-----+
|      3 | SCIT    |     210 |   20 |
|      3 | SCIT    |     211 |   22 |
+-----+-----+-----+-----+
```

ROOM table

```
+-----+-----+-----+-----+
| bldgnum | bldgname | roomnum | area |
+-----+-----+-----+-----+
|      3 | SCIT    |     210 |   20 |
|      3 | SCIT    |     211 |   22 |
|      3 | SCIT    |     213 |   20 |
+-----+-----+-----+-----+
```

ROOM table

# Insertion test

ROOM table

bldgnum	bldgname	roomnum	area
3	SCIT	210	20
3	SCIT	211	22
3	SCIT	213	20

## Problems:

- It clearly visible that **building name** is repeated as many times as many rooms are included in a building

# Insertion test

In another example, we insert few rows into a relational table

WAREHOUSE

```
WAREHOUSE(name, address, part, quantity)
```

```
PRIMARY KEY = (name, part)
```

WAREHOUSE schema

```
+-----+-----+-----+-----+
| name      | address      | part  | quantity |
+-----+-----+-----+-----+
| Golden Bolts | Northfields Ave | bolt  | 210 |
+-----+-----+-----+-----+
```

WAREHOUSE table

```
+-----+-----+-----+-----+
| name      | address      | part  | quantity |
+-----+-----+-----+-----+
| Golden Bolts | Northfields Ave | bolt  | 210 |
| Golden Bolts | Northfields Ave | lock  | 20 |
+-----+-----+-----+-----+
```

WAREHOUSE table

```
+-----+-----+-----+-----+
| name      | address      | part  | quantity |
+-----+-----+-----+-----+
| Golden Bolts | Northfields Ave | bolt  | 210 |
| Golden Bolts | Northfields Ave | lock  | 20 |
| Golden Bolts | Northfields Ave | screw | 211 |
+-----+-----+-----+-----+
```

WAREHOUSE table



# Insertion test

WAREHOUSE table

name	address	part	quantity
Golden Bolts	Northfields Ave	bolt	210
Golden Bolts	Northfields Ave	lock	20
Golden Bolts	Northfields Ave	screw	211

## Problems:

- **Address** of a warehouse is repeated as many time as many different parts are stored in the warehouse
- If at some point in time there are no parts stored in a warehouse then there may be no rows to keep a warehouse **address** or the values of certain attributes must be set to **NULL**

# Insertion test

In yet another example, we insert few rows into a relational table

EMPLOYEE

```
EMPLOYEE(enum, skill, hobby)
PRIMARY KEY = (enum, skill, hobby)
```

EMPLOYEE schema

```
+-----+-----+-----+
| enum | skill   | hobby   |
+-----+-----+-----+
| 7    | cooking | hiking  |
+-----+-----+-----+
```

EMPLOYEE table

```
+-----+-----+-----+
| enum | skill   | hobby   |
+-----+-----+-----+
| 7    | cooking | hiking  |
| 7    | cooking | swimming|
+-----+-----+-----+
```

EMPLOYEE table

```
+-----+-----+-----+
| enum | skill   | hobby   |
+-----+-----+-----+
| 7    | cooking | hiking  |
| 7    | cooking | swimming|
| 7    | programming | hiking |
| 7    | programming | swimming|
+-----+-----+-----+
```

EMPLOYEE table

# Insertion test

enum	skill	hobby
7	cooking	hiking
7	cooking	swimming
7	programming	hiking
7	programming	swimming

EMPLOYEE table

## Problems:

- Skill name must be repeated with each hobby name
- Hobby name must be repeated with each skill name
- If at some point in time an employee has no hobbies (or skills) then a value of attribute hobby (or skill) must be set to NULL, however, it is impossible due to PRIMARY KEY = (enum, skill, hobby) constraint

# Database Design Quality

## Outline

Why not ONE BIG TABLE !?

Where is a problem ?

Insertion test

Join test

Deletion test

Update test

Good design guidelines

# Join test

We consider the relational tables with the following schemas:

SUPPLIER(sname, city)  
PRIMARY KEY = (sname, city)

SUPPLIER schema

COMPANY(cname, city)  
PRIMARY KEY = (cname, city)

COMPANY schema

```
+-----+-----+
| sname | city |
+-----+-----+
| Harry | Paris |
| James | Paris |
| Robin | Rome  |
+-----+-----+
```

SUPPLIER table

```
+-----+-----+
| cname      | city |
+-----+-----+
| Golden Bolts | Paris |
| Golden Bolts | Rome  |
| Lazy Lobster | Rome  |
+-----+-----+
```

COMPANY table

# Join test

The result of join of `SUPPLIER` and `COMPANY` tables over a column `city`:

sname	city	cname	city
Harry	Paris	Golden Bolts	Paris
James	Paris	Golden Bolts	Paris
Robin	Rome	Golden Bolts	Rome
Robin	Rome	Lazy Lobster	Rome

JOIN of SUPPLIER and COMPANY tables

## Problems:

- Join of relational tables `SUPPLIER` and `COMPANY` creates the spurious row `[Robin | Rome | Lazy Lobster | Rome]` that represent wrong information

# Database Design Quality

## Outline

Why not ONE BIG TABLE !?

Where is a problem ?

Insertion test

Join test

Deletion test

Update test

Good design guidelines

# Deletion test

We consider a relational table with the following schema:

```
SHIPMENT(sname, product, part)
PRIMARY KEY = (sname, product, part)
```

SHIPMENT schema

- The table contains the following rows:

```
+-----+-----+-----+
| sname | product | part   |
+-----+-----+-----+
| James | audio system | amplifier |
| James | audio system | speakers  |
| James | computer   | hard disk |
| James | computer   | mainboard |
| James | computer   | processor |
+-----+-----+-----+
```

SHIPMENT table

- Deletion of a product `computer` requires deletion of additional two rows
- Then, deletion of a product `audio system` causes accidental deletion of information about supplier `James`

Problems:

- Deletion of a row triggers deletion of the other rows in the same table
- Deletion of a row accidentally deletes other information



# Database Design Quality

## Outline

Why not ONE BIG TABLE !?

Where is a problem ?

Insertion test

Join test

Deletion test

Update test

Good design guidelines

# Update test

We consider a relational table with the following schema:

```
WAREHOUSE(name, address, part, quantity)
PRIMARY KEY = (name, part)
```

WAREHOUSE schema

- The table contains the following rows:

name	address	part	quantity
Golden Bolts	Northfields Ave	bolt	210
Golden Bolts	Northfields Ave	lock	20
Golden Bolts	Northfileds Ave	screw	211

WAREHOUSE table

- Modification of an address `Northfields Ave` requires replication of a modification in two other rows

Problem:

- Modification of a row triggers modifications of the other rows in the same table

# Database Design Quality

## Outline

Why not ONE BIG TABLE !?

Where is a problem ?

Insertion test

Join test

Deletion test

Update test

Good design guidelines

# Good design guidelines

Design a relational table such that it is easy to explain its meaning

Do not combine unrelated attributes into the same table

Design a relational table such that insertion, deletion and update tests do not cause problems

Minimize the number of attributes whose values can be missing ([NULL](#))

Design the relational tables such that they can be joined with equality conditions on attributes that are either primary or foreign keys in a way that no spurious rows are generated

**And the first of all ALWAYS START FROM CONCEPTUAL MODELING !**

# References

T. Connolly, C. Begg, Database Systems, A Practical Approach to Design, Implementation, and Management, Chapters 14.1 - 14.3 Introduction to normalization, Pearson Education Ltd, 2015