

SELECT Statement (3)

Dr Janusz R. Getta

School of Computing and Information Technology -
University of Wollongong

SELECT statement (3)

Outline

Join queries

Natural join queries

Column name join queries

Cross join queries

Join queries over more than 2 tables

Self-join queries

Join queries

Join operation "connects" the rows from two relational tables

```
SELECT *  
FROM TABLE_1 T1 JOIN TABLE_2 T2  
ON Condition;
```

TABLE_1 T1

COLUMN_1	COLUMN_2	...	COLUMN_N
		...	
		...	
		...	

TABLE_2 T2

COLUMN_1	COLUMN_2	...	COLUMN_N
		...	
		...	
		...	



RESULTS

T1.COLUMN_1	T1.COLUMN_2	...	T1.COLUMN_N	T2.COLUMN_1	T2.COLUMN_2	...	T1.COLUMN_N
		
		

Join queries

Join operation "connects" the rows from two relational tables


```
SELECT *  
FROM TABLE_1 T1 JOIN TABLE_2 T2  
ON Condition;
```

TABLE_1 T1

COLUMN_1	COLUMN_2	...	COLUMN_N
		...	
		...	
		...	

TABLE_2 T2

COLUMN_1	COLUMN_2	...	COLUMN_N
		...	
		...	
		...	



RESULTS

T1.COLUMN_1	T1.COLUMN_2	...	T1.COLUMN_N	T2.COLUMN_1	T2.COLUMN_2	...	T1.COLUMN_N
		
		

Join queries

Join operation "connects" the rows from two relational tables


```
SELECT *  
FROM TABLE_1 T1 JOIN TABLE_2 T2  
ON Condition;
```

TABLE_1 T1

COLUMN_1	COLUMN_2	...	COLUMN_N
		...	
		...	
		...	

TABLE_2 T2

COLUMN_1	COLUMN_2	...	COLUMN_N
		...	
		...	
		...	



RESULTS

T1.COLUMN_1	T1.COLUMN_2	...	T1.COLUMN_N	T2.COLUMN_1	T2.COLUMN_2	...	T1.COLUMN_N
		
		
		
		

Join queries

Sample database

```
CREATE TABLE DEPARTMENT(  
  name          VARCHAR(50)          NOT NULL,  
  code          CHAR(5)              NOT NULL,  
  total_staff_number DECIMAL(2)      NOT NULL,  
  chair         VARCHAR(50)          NULL,  
  budget        DECIMAL(9,1)         NOT NULL,  
  CONSTRAINT dept_pkey PRIMARY KEY(name),  
  CONSTRAINT dept_cke1 UNIQUE(code),  
  CONSTRAINT dept_cke2 UNIQUE(chair),  
  CONSTRAINT dept_check1 CHECK (total_staff_number BETWEEN 1 AND 50) );
```

CREATE TABLE statement

```
CREATE TABLE COURSE(  
  cnum          CHAR(7)              NOT NULL,  
  title         VARCHAR(200)         NOT NULL,  
  credits       DECIMAL(2)           NOT NULL,  
  offered_by    VARCHAR(50)          NULL,  
  CONSTRAINT course_pkey PRIMARY KEY(cnum),  
  CONSTRAINT course_check1 CHECK (credits IN (6, 12)),  
  CONSTRAINT course_fkey1 FOREIGN KEY(offered_by)  
    REFERENCES DEPARTMENT(name) ON DELETE CASCADE );
```

CREATE TABLE statement

Join queries

Consider the following query: Find the titles of all courses offered by a department chaired by Peter

DEPARTMENT

name | code | total_staff_number | chair | budget

Relational schema

- There are no titles of courses in a relational table DEPARTMENT !
- The titles of courses are in a relational table COURSE

COURSE

cnum | title | credits | offered_by

Relational schema

- To implement the query we must use two tables: DEPARTMENT and COURSE
- The rows from a table DEPARTMENT must be joined with (connected to) the respective rows in a relational table COURSE over a condition
`DEPARTMENT.name = COURSE.offered_by`

SELECT statement with JOIN operation

```
SELECT COURSE.title
FROM COURSE JOIN DEPARTMENT
      ON DEPARTMENT.name = COURSE.offered_by
WHERE DEPARTMENT.chair = 'Peter';
```

Join queries

Implementation of the query **Find the titles of all courses offered by a department chaired by Peter** has the following syntactical variations

SELECT statement with JOIN operation

```
SELECT title
FROM COURSE JOIN DEPARTMENT
      ON name = offered_by
WHERE chair = 'Peter';
```

SELECT statement with JOIN operation

```
SELECT C.title
FROM COURSE C JOIN DEPARTMENT D
      ON D.name = C.offered_by
WHERE D.chair = 'Peter';
```

SELECT statement with JOIN operation

```
SELECT COURSE.title
FROM COURSE, DEPARTMENT
WHERE DEPARTMENT.name = COURSE.offered_by AND DEPARTMENT.chair = 'Peter';
```

SELECT statement with JOIN operation

```
SELECT title
FROM COURSE, DEPARTMENT
WHERE name = offered_by AND chair = 'Peter';
```


ANSI SQL Syntax

The following implementations of the query **Find the titles of all courses offered by a department chaired by Peter** are consistent with ANSI SQL standard

SELECT statement with JOIN operation (ANSI SQL standard)

```
SELECT COURSE.title
FROM COURSE JOIN DEPARTMENT
      ON DEPARTMENT.name = COURSE.offered_by
WHERE DEPARTMENT.chair = 'Peter';
```

SELECT statement with JOIN operation (ANSI SQL standard)

```
SELECT title
FROM COURSE JOIN DEPARTMENT
      ON name = offered_by
WHERE chair = 'Peter';
```

SELECT statement with JOIN operation (ANSI SQL standard)

```
SELECT C.title
FROM COURSE C JOIN DEPARTMENT D
      ON D.name = C.offered_by
WHERE D.chair = 'Peter';
```

SELECT statement (3)

Outline

Join queries

Natural join queries

Column name join queries

Cross join queries

Join queries over more than 2 tables

Self-join queries

Natural join queries

Consider a query: Find the names of all employees from a department chaired by James Bond over the relational tables

DEPARTMENT

Relational schema

dname | code | total staff number | chair | budget

EMPLOYEE

Relational schema

enum | ename | dname

A natural join query

SELECT statement with NATURAL JOIN operation

```
SELECT ename
FROM EMPLOYEE NATURAL JOIN DEPARTMENT
WHERE chair = 'James Bond';
```

- is equivalent to a join query

SELECT statement with JOIN operation (equivalent to a statment above)

```
SELECT ename
FROM EMPLOYEE JOIN DEPARTMENT
      ON EMPLOYEE.dname = DEPARTMENT.dname
WHERE chair = 'James Bond';
```

SELECT statement (3)

Outline

Join queries

Natural join queries

Column name join queries

Cross join queries

Join queries over more than 2 tables

Self-join queries

Column name join queries

Consider a query: Find the names of all employees from a department chaired by James Bond over the relational tables

DEPARTMENT

Relational schema

dname | code | total staff number | chair | budget

EMPLOYEE

Relational schema

enum | ename | dname

A column name join query

SELECT statement with column name JOIN operation

```
SELECT ename
FROM EMPLOYEE JOIN DEPARTMENT
      USING(dname)
WHERE chair = 'James Bond';
```

- is equivalent to a join query

SELECT statement with JOIN operation (equivalent to a statement above)

```
SELECT ename
FROM EMPLOYEE JOIN DEPARTMENT
      ON EMPLOYEE.dname = DEPARTMENT.dname
WHERE chair = 'James Bond';
```

SELECT statement (3)

Outline

Join queries

Natural join queries

Column name join queries

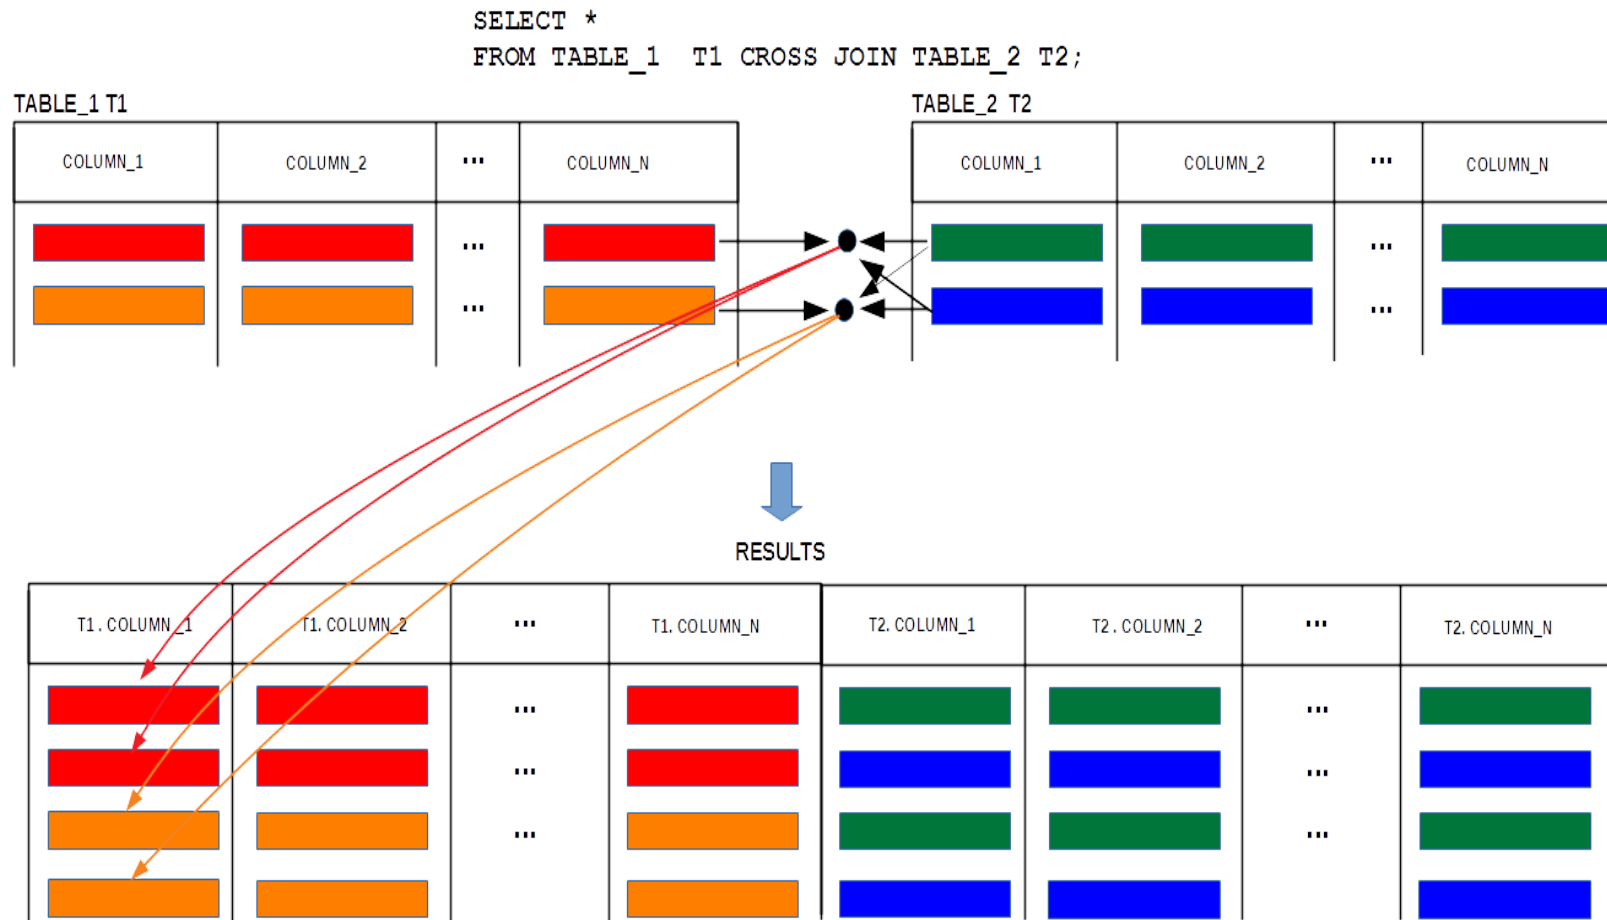
Cross join queries

Join queries over more than 2 tables

Self-join queries

Cross join queries

Cross join operation "connects" all rows from a relational table with all rows from another relational table



Cross join queries

Consider a query: Find all pairs of the names of employees and the names of chair people over the relational tables

DEPARTMENT

Relational schema

dname | code | total staff number | chair | budget

EMPLOYEE

Relational schema

enum | ename | dname

A cross join query

```
SELECT ename, chair
FROM EMPLOYEE CROSS JOIN DEPARTMENT;
```

SELECT statement with CROSS JOIN operations

- is equivalent to the following join queries

```
SELECT ename, chair
FROM EMPLOYEE JOIN DEPARTMENT;
```

SELECT statement equivalent to a statement above

```
SELECT ename, chair
FROM EMPLOYEE, DEPARTMENT;
```

SELECT statement equivalent to a statement above

SELECT statement (3)

Outline

Join queries

Natural join queries

Column name join queries

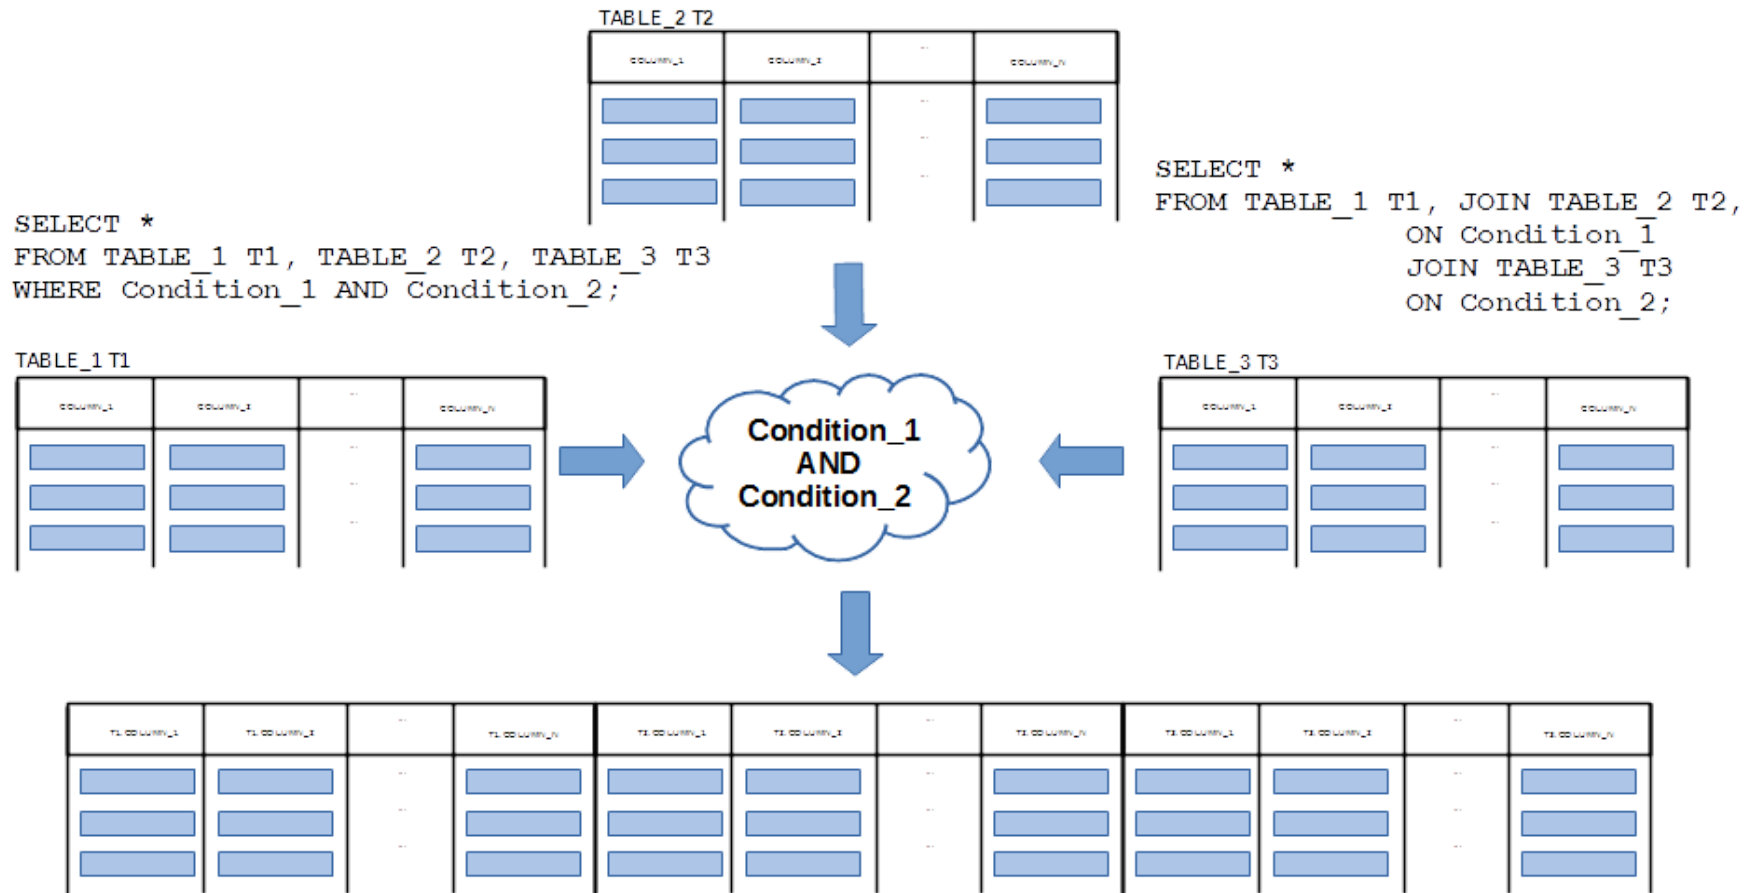
Cross join queries

Join queries over more than 2 tables

Self-join queries

Join queries over more than 2 tables

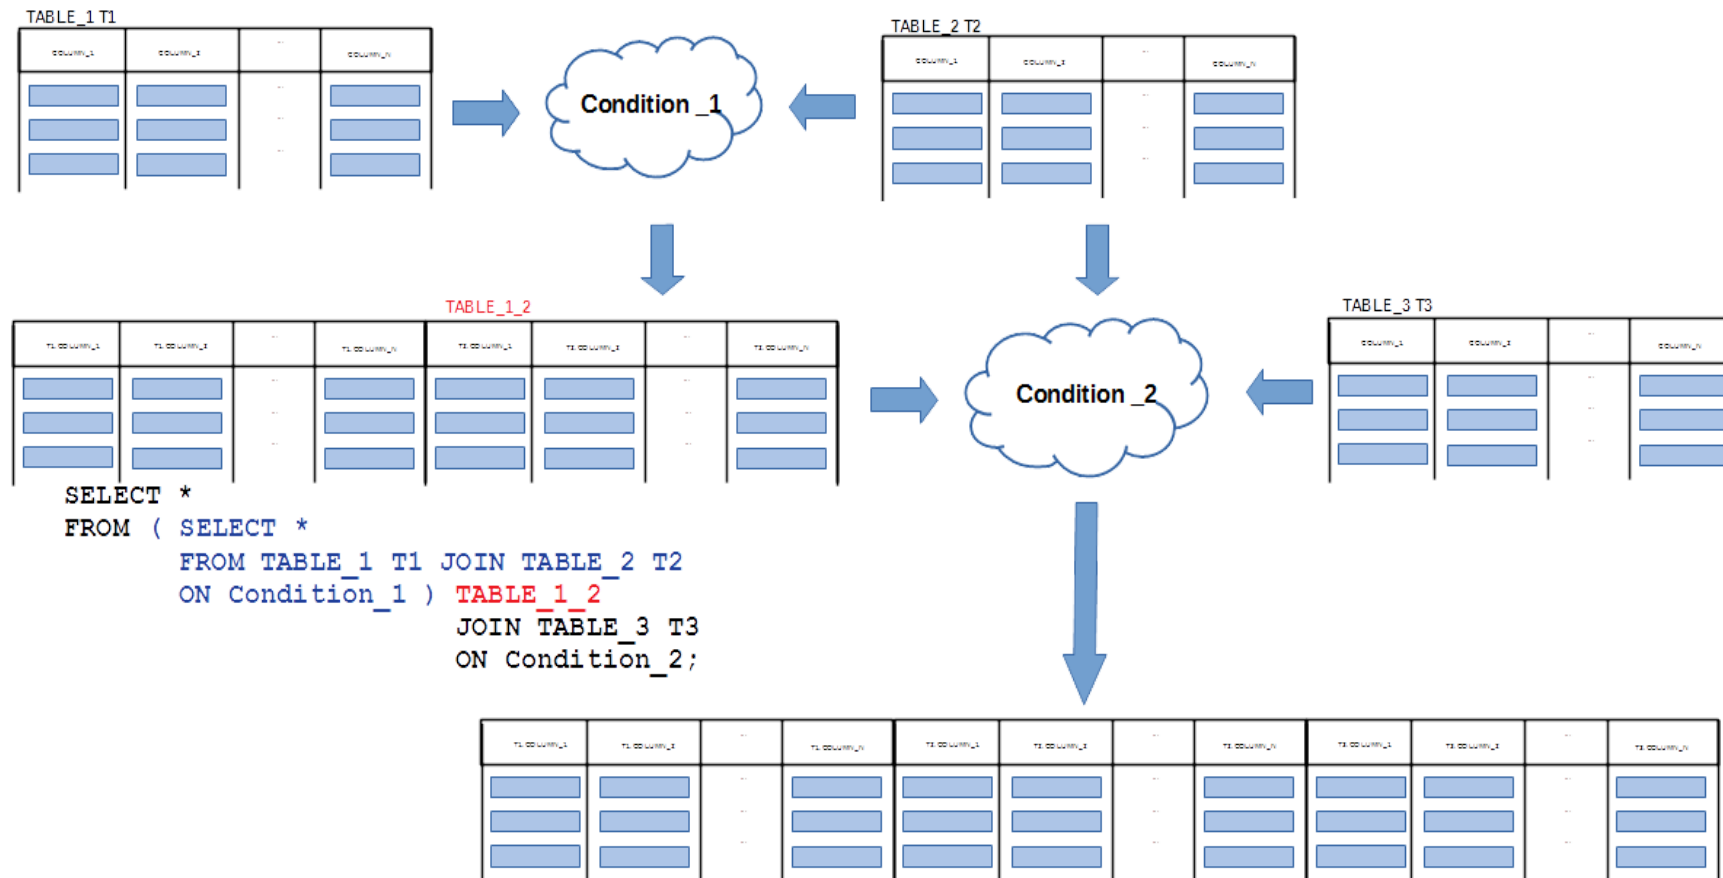
A sample join of three relational tables `TABLE_1`, `TABLE_2`, and `TABLE_3`



Join queries over more than 2 tables

Application of **inline view** to simplify join of three relational tables

TABLE_1, TABLE_2, and TABLE_3



Join queries over more than 2 tables

Consider the relational tables with the following schemas

COURSE

cnum | title | credits

Relational schema

STUDENT

snum | name | degree

Relational schema

ENROLMENT

cnum | snum | edate | result

Relational schema

A query **find the names of all students who enrolled java course** can be implemented as the following join query

SELECT statement that joins three relational tables

```
SELECT STUDENT.name
FROM COURSE JOIN ENROLMENT
      ON COURSE.cnum = ENROLMENT.cnum
      JOIN STUDENT
      ON ENROLMENT.snum = STUDENT.snum
WHERE COURSE.title = 'Java';
```

Join queries over more than 2 tables

Implementation of a query **find the names of all students who enrolled Java course** has the following syntactical variations

```
SELECT STUDENT.name
FROM COURSE JOIN ENROLMENT
      ON COURSE.cnum = ENROLMENT.cnum
      JOIN STUDENT
      ON ENROLMENT.snum = STUDENT.snum
WHERE COURSE.title = 'Java';
```

SELECT statement that joins three relational tables

```
SELECT STUDENT.name
FROM COURSE, ENROLMENT, STUDENT
WHERE COURSE.cnum = ENROLMENT.cnum AND ENROLMENT.snum = STUDENT.snum AND
      COURSE.title = 'Java';
```

SELECT statement that joins three relational tables

```
SELECT STUDENT.name
FROM ( SELECT *
      FROM COURSE JOIN ENROLMENT
            ON COURSE.cnum = ENROLMENT.cnum
      WHERE COURSE.title = 'Java' ) CE JOIN STUDENT
      ON CE.snum = STUDENT.snum
WHERE COURSE.title = 'Java';
```

SELECT statement that joins three relational tables

SELECT statement (3)

Outline

Join queries

Natural join queries

Column name join queries

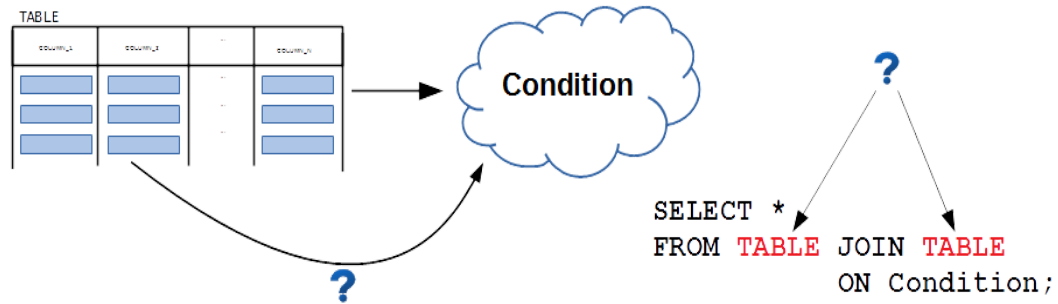
Cross join queries

Join queries over more than 2 tables

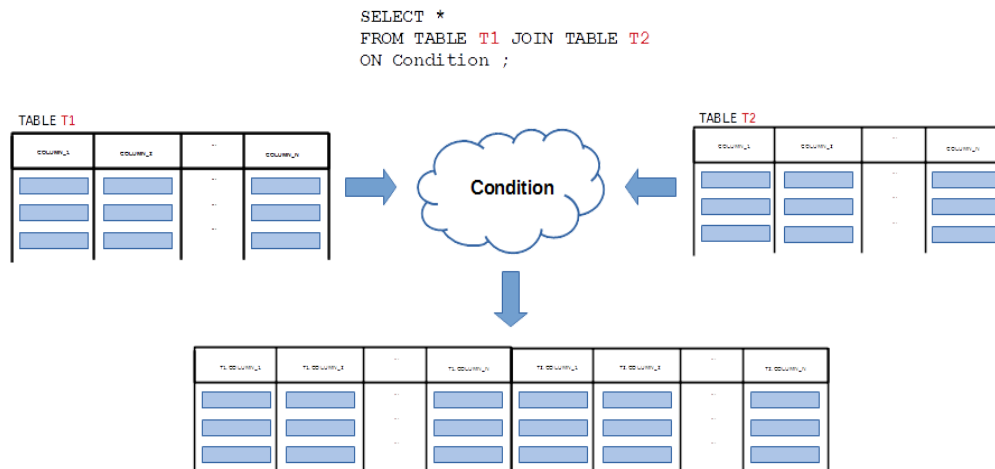
Self-join queries

Self-join queries

What if a relational table must be joined with itself ?



A table that must be joined with itself obtains two different **alias names**



Self-join queries

Consider a relational table with the following schema

EMPLOYEE
enum, name, manager

Relational schema

and the following contents

enum	name	manager
10	John	NULL
20	Peter	10
30	Mary	10
40	Mike	20
50	Kate	20
60	Greg	50
70	Phil	50

Relational table

Consider a query **find a name of manager of employee number 40**

We can "plan" the implementation in the following way

- (1) Find an employee number of manager of employee number 40
- (2) Find a name of employee found in the previous query

Self-join queries

Implementation of a plan

- (1) Find employee number of manager of employee number 40
- (2) Find a name of employee found in the previous query

is the following

<pre>SELECT manager FROM EMPLOYEE WHERE enum = 40;</pre>	SELECT statement
20	Result
<pre>SELECT name FROM EMPLOYEE WHERE enum = 20;</pre>	SELECT statement
Peter	Result

Is it possible to implement the query as one `SELECT` statement ?

YES ! In more than one way !

Self-join queries

Solution 1

Assume that we have two identical relational tables `E1` and `E2`

E1				E2				Relational tables	
enum	name	manager		enum	name	manager			
10	John	NULL		10	John	NULL			
20	Peter	10	----->	20	Peter	10			
30	Mary	10		30	Mary	10			
40	Mike	20	-----	40	Mike	20			
50	Kate	20		50	Kate	20			
60	Greg	50		60	Greg	50			
70	Phil	50		70	Phil	50			

To find a name of manager of employee number 40 we take a row

40	Mike	20		Row
----	------	----	--	-----

- from a table `E1` and we join over a condition `E1.manager = E2.enum` with a row

20	Peter	10		Row
----	-------	----	--	-----

- from a table `E2`

Self-join queries

So, how do we implement such "magic" ?

```
SELECT E2.name
FROM EMPLOYEE E1 JOIN EMPLOYEE E2
      ON E1.manager = E2.enum
WHERE E1.enum = 40;
```

SELECT statement with self-join

E1				E2			
enum	name	manager		enum	name	manager	
10	John	NULL		10	John	NULL	
20	Peter	10	----->	20	Peter	10	
30	Mary	10		30	Mary	10	
40	Mike	20	-----	40	Mike	20	
50	Kate	20		50	Kate	20	
60	Greg	50		60	Greg	50	
70	Phil	50		70	Phil	50	

Relational tables

Self-join queries

Solution 2

We use **inline views** technique to combine the following queries

```
SELECT manager
FROM EMPLOYEE
WHERE enum = 40;
```

SELECT statement

```
SELECT name
FROM EMPLOYEE
WHERE enum = 20;
```

SELECT statement

We use the first **SELECT** statement to create an **inline view** E40

```
( SELECT manager
  FROM EMPLOYEE
 WHERE enum = 40 ) E40
```

Inline view

Then we join an **inline view** E40 with a relational table **EMPLOYEE**

```
SELECT EMPLOYEE.name
FROM EMPLOYEE JOIN ( SELECT manager
                     FROM EMPLOYEE
                     WHERE enum = 40 ) E40
ON EMPLOYEE.enum = E40.manager;
```

SELECT statement with inline view

Self-join queries

In another query we find the names of all employees directly managed by Kate

enum	name	manager
10	John	NULL
20	Peter	10
30	Mary	10
40	Mike	20
50	Kate	20
60	Greg	50
70	Phil	50

Relational table

We "plan" the implementation in the following way

- (1) Find an employee number of an employee Kate
- (2) Find the names of employees who have a number found in the previous query in a column `manager`

Self-join queries

Solution 1

Assume that we have two identical relational tables `E1` and `E2`

E1				E2				Relational tables			
enum	name	manager		enum	name	manager					
10	John	NULL		10	John	NULL					
20	Peter	10		20	Peter	10					
30	Mary	10		30	Mary	10					
40	Mike	20		40	Mike	20					
50	Kate	20		50	Kate	20					
60	Greg	50	<---	60	Greg	50					
70	Phil	50	<---	70	Phil	50					

To find a number of employee Kate we take a row

50	Kate	20		Row
----	------	----	--	-----

- from a table `E2` and we join it over a condition `E2.enum = E1.manager` with the rows

60	Greg	50		Row
----	------	----	--	-----

70	Phil	50		Row
----	------	----	--	-----

- from a table `E1`

Self-join queries

So, how do we implement such "magic" ?

```
SELECT E1.name
FROM EMPLOYEE E1 JOIN EMPLOYEE E2
      ON E1.manager = E2.enum
WHERE E2.name = 'Kate';
```

SELECT statement with self-join

E1	E2
enum name manager	enum name manager
10 John NULL	10 John NULL
20 Peter 10	20 Peter 10
30 Mary 10	30 Mary 10
40 Mike 20	40 Mike 20
50 Kate 20	50 Kate 20
60 Greg 50 <---	60 Greg 50
70 Phil 50 <---	70 Phil 50

Relational tables

Self-join queries

Solution 2

We use **inline views** technique to combine the following queries

```
SELECT enum
FROM EMPLOYEE
WHERE name = 'Kate';
```

SELECT statement

```
SELECT name
FROM EMPLOYEE
WHERE manager = 50;
```

SELECT statement

We use the first **SELECT** statement to create an **inline view** **KATE**

```
( SELECT enum
  FROM EMPLOYEE
  WHERE name = 'Kate' ) KATE
```

Inline view

Then we join an **inline view** **KATE** with a relational table **EMPLOYEE**

```
SELECT EMPLOYEE.name
FROM EMPLOYEE JOIN ( SELECT enum
                     FROM EMPLOYEE
                     WHERE name = 'Kate' ) KATE
ON EMPLOYEE.manager = KATE.enum;
```

SELECT statement with inline view

References

T. Connolly, C. Begg, Database Systems, A Practical Approach to Design, Implementation, and Management, Chapters 6.3.7 Multi-table Queries, Pearson Education Ltd, 2015

D. Darmawikarta, SQL for MySQL A Beginner's Tutorial, Chapter 6, pages 55 - 61 Brainy Software Inc. First Edition: June 2014

[How to ... ? Cookbook, How to implement queries in SQL ? \(Part 1\), Recipe 5.4 How to implement simple join queries ?](#)

[How to ... ? Cookbook, How to implement queries in SQL ? \(Part 2\) Recipe 6.1 How to implement self join queries ?](#)