

SELECT Statement (5)

Dr Janusz R. Getta

School of Computing and Information Technology -
University of Wollongong

SELECT statement (5)

Outline

Nested queries

Correlated nested queries

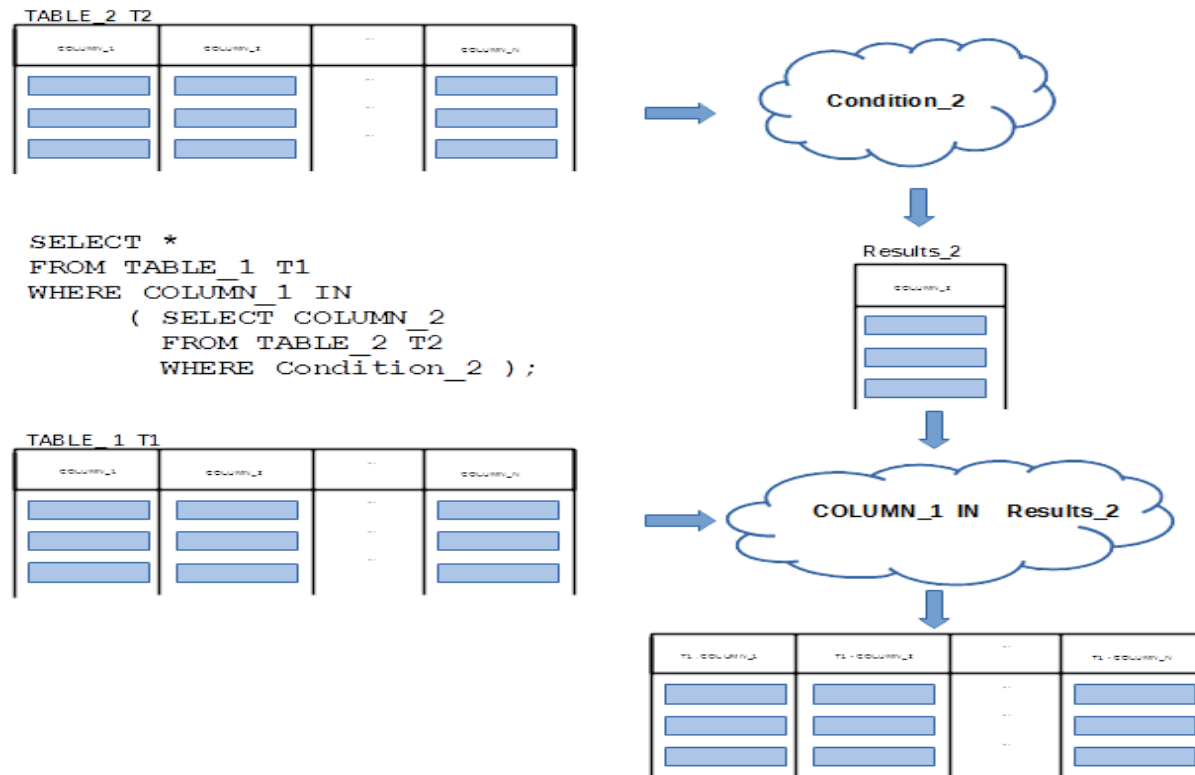
Queries with WITH clause

Relational views

Advanced DML statements

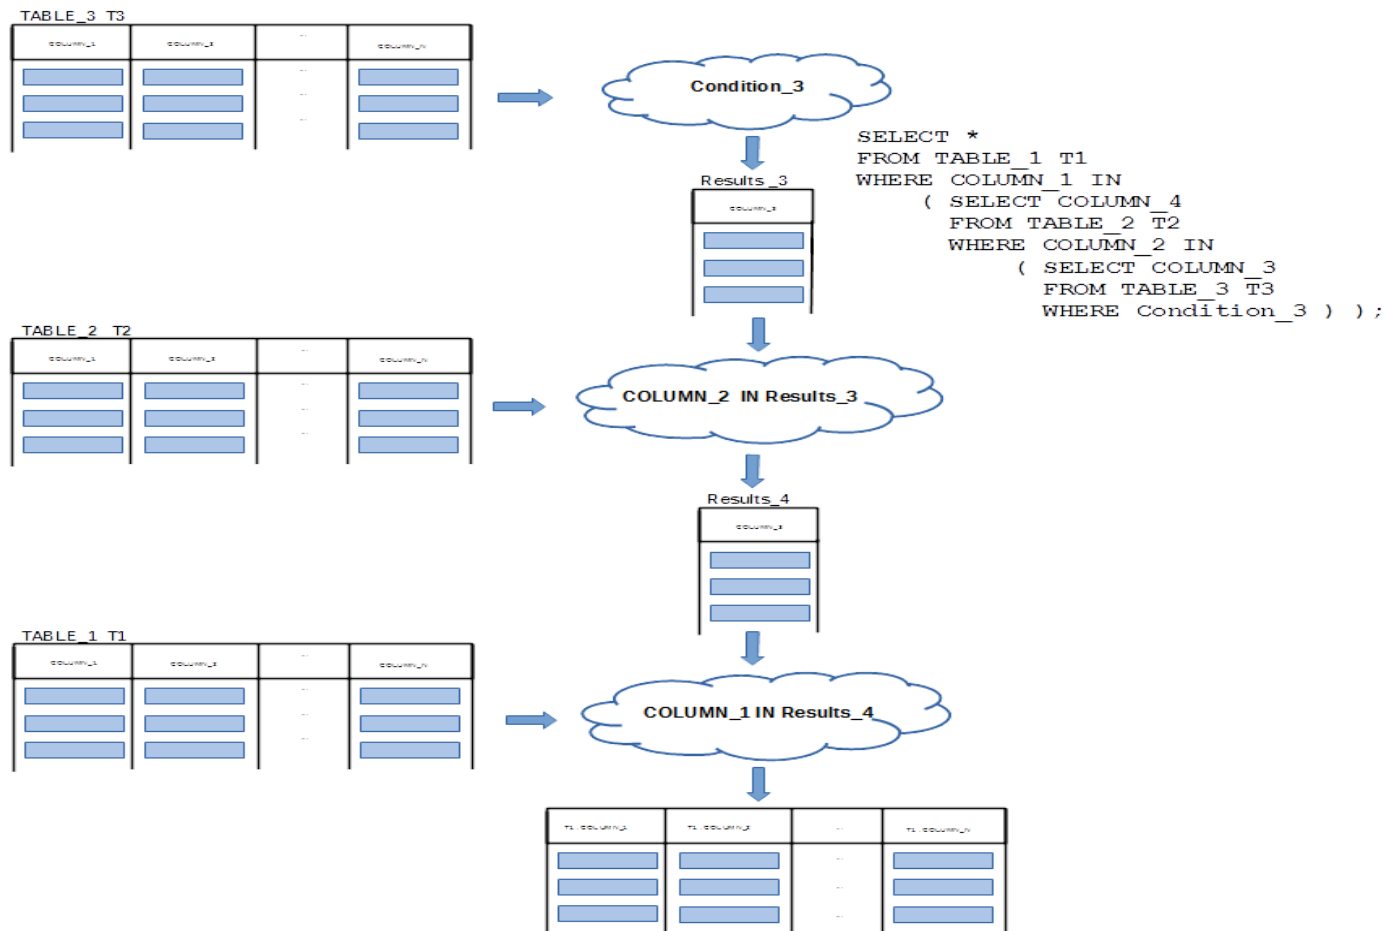
Nested queries

Nested query is a query with another query embedded in **WHERE**, **SELECT**, **FROM** clauses



Nested queries

Nested query is a query with another query embedded in **WHERE**, **SELECT**, **FROM** clauses



Nested queries

Sample database

```
CREATE TABLE DEPARTMENT(  
  name          VARCHAR(50)          NOT NULL,  
  code          CHAR(5)              NOT NULL,  
  total_staff_number DECIMAL(2)      NOT NULL,  
  chair         VARCHAR(50)          NULL,  
  budget        DECIMAL(9,1)         NOT NULL,  
  CONSTRAINT dept_pkey PRIMARY KEY(name),  
  CONSTRAINT dept_cke1 UNIQUE(code),  
  CONSTRAINT dept_cke2 UNIQUE(chair),  
  CONSTRAINT dept_check1 CHECK (total_staff_number BETWEEN 1 AND 50) );
```

CREATE TABLE statement

```
CREATE TABLE COURSE(  
  cnum          CHAR(7)              NOT NULL,  
  title         VARCHAR(200)         NOT NULL,  
  credits       DECIMAL(2)           NOT NULL,  
  offered_by    VARCHAR(50)          NULL,  
  CONSTRAINT course_pkey PRIMARY KEY(cnum),  
  CONSTRAINT course_check1 CHECK (credits IN (6, 12)),  
  CONSTRAINT course_fkey1 FOREIGN KEY(offered_by)  
    REFERENCES DEPARTMENT(name) ON DELETE CASCADE );
```

CREATE TABLE statement

Nested queries

A class of **nested queries** is based on a concept of **inline views**

Inline views can be used to reduce the complexity of query implementation

An **inner query** is created first and its outcomes are used as a relational table (**inline view**) in an **outer query**

For example, a query **find the titles of courses offered by a department chaired by Peter** is decomposed into the following two queries:

- Q1: Find a department chaired by Peter
- Q2: Find the titles of courses offered by a department found in Q1

A query Q1 is implemented first and then it is used in **WHERE** clause of query Q2

Nested queries

In **nested queries** **SELECT** statements are nested to theoretically unlimited level in **WHERE** clause

```
SELECT title
FROM COURSE
WHERE offered_by IN ( SELECT name
                      FROM DEPARTMENT
                      WHERE chair = 'Peter' );
```

Nested query with a set membership operation IN

If inner query returns more than one row the we must use **IN** instead of **=**

Nested queries

A way how we implement a query **find the titles of courses offered by a department chaired by Peter** is the following

Create a inner query as **inline view** Q

```
( SELECT name  
  FROM DEPARTMENT  
 WHERE chair = 'Peter' ) Q
```

Inline view

Create outer query that references an **inline view** Q created earlier

```
SELECT title  
FROM COURSE  
WHERE offered_by IN Q.name
```

Query that references an inline view Q

Replace a reference Q to an **inline view** with the **inline view**

```
SELECT title  
FROM COURSE  
WHERE offered_by IN ( SELECT name  
                      FROM DEPARTMENT  
                      WHERE chair = 'Peter' );
```

Nested query with a set membership operation IN

Nested queries

Another example, find the chairs of all departments that offer 12 credit point courses

An inner query as inline view Q

```
( SELECT offered_by  
  FROM COURSE  
 WHERE credits = 12 ) Q;
```

Inline view

Create outer query that references an inline view Q created earlier

```
SELECT chair  
FROM DEPARTMENT  
WHERE name IN Q.offered_by
```

Query that references an inline view Q

Replace a reference to an inline view Q with the inline view itself

```
SELECT chair  
FROM DEPARTMENT  
WHERE name IN ( SELECT offered_by  
                 FROM COURSE  
                 WHERE credits = 12 );
```

Nested query with a set membership operation IN

Nested queries

A query **find the chairs of all departments that offer no courses** is an example of **ANTIJOIN** operation

It is equivalent to a query **find all rows in a relational table DEPARTMENT that cannot be joined with any row in a relational table COURSE**

An inner query as **inline view** finds all departments that offer at least one course

```
( SELECT offered_by  
  FROM COURSE ) Q;
```

Inline view

An outer query **finds all chairs of departments that are NOT included in the results of an inner query**

```
SELECT chair  
FROM DEPARTMENT  
WHERE name NOT IN Q.offered_by
```

Query that references an inline view

Nested queries

Finally, we replace a reference to an **inline view** Q with the **inline view** itself

```
SELECT chair
FROM DEPARTMENT
WHERE name NOT IN ( SELECT offered_by
                    FROM COURSE );
```

Nested query with a negated set membership operation IN

SELECT statement (5)

Outline

Nested queries

Correlated nested queries

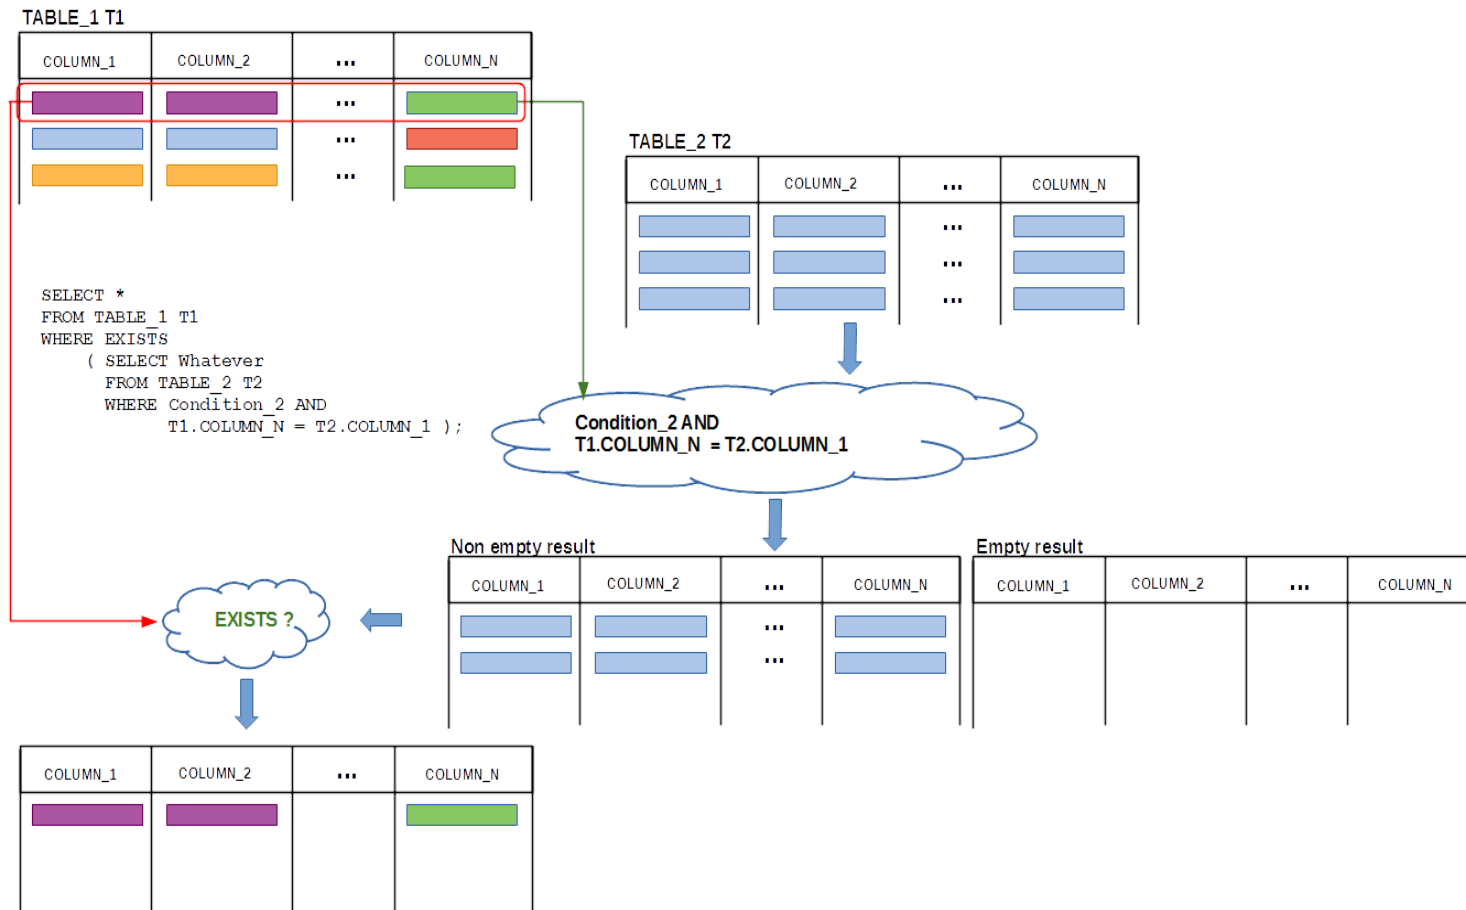
Queries with WITH clause

Relational views

Advanced DML statements

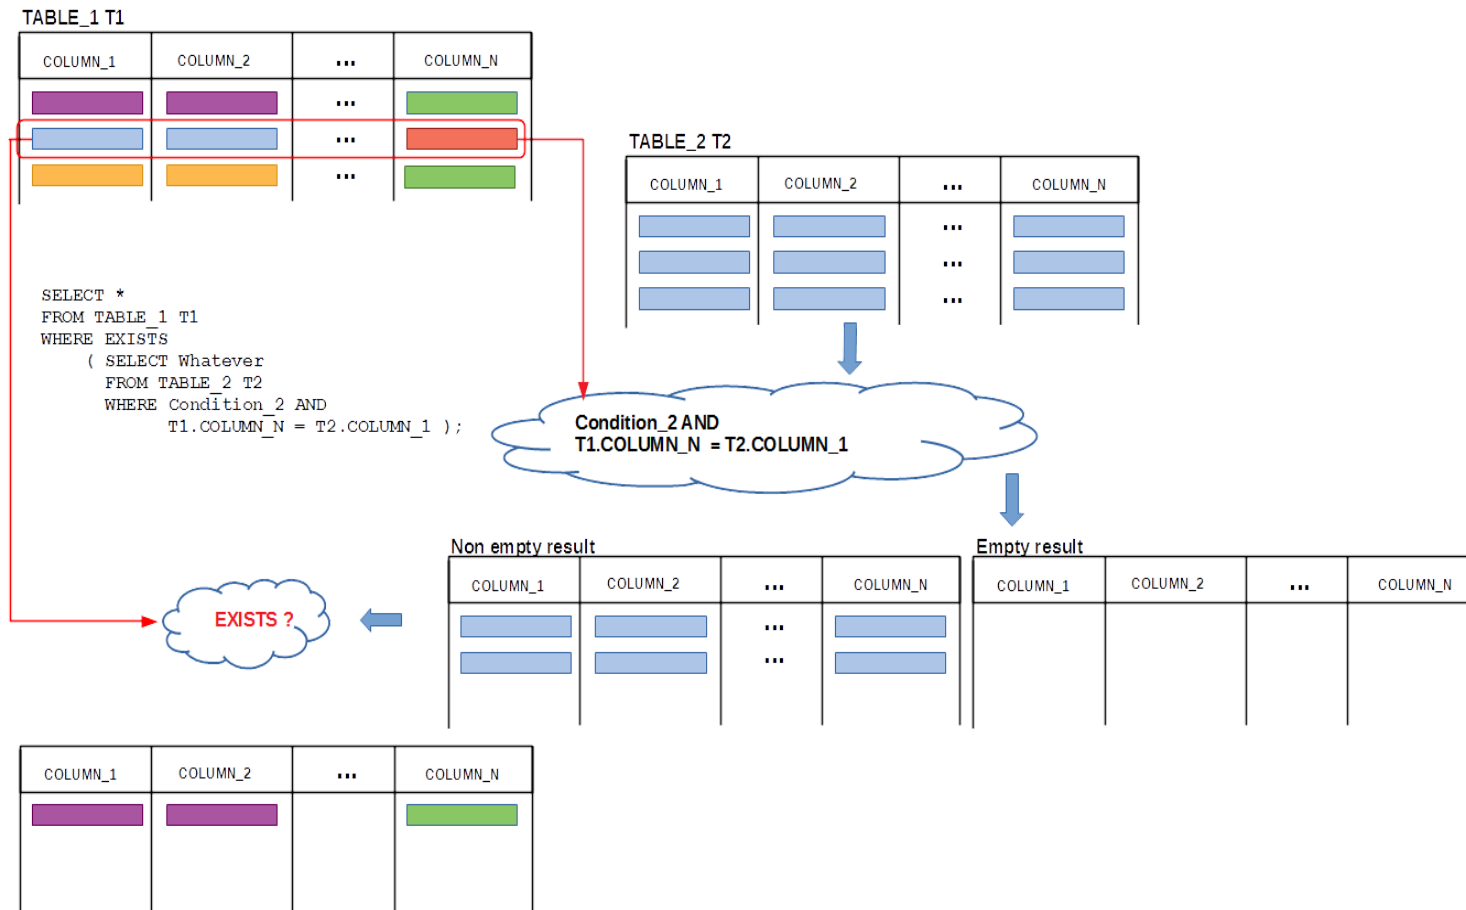
Correlated nested queries

In a **correlated nested query** an **inline view** may reference the names of relational tables used in **SELECT** statement outer to the **inline view**



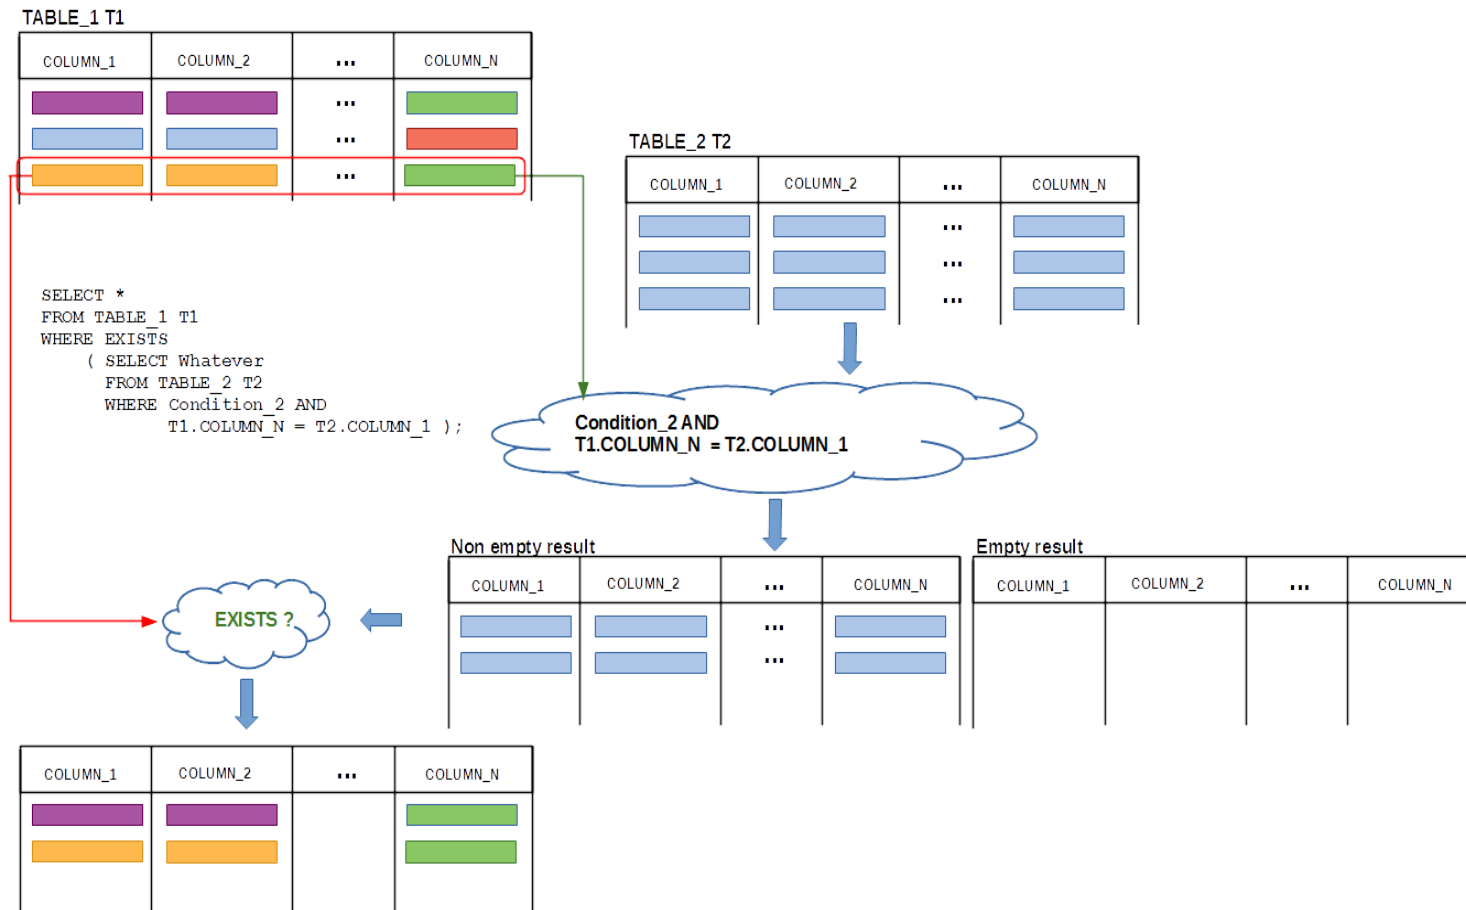
Correlated nested queries

In a **correlated nested query** an **inline view** may reference the names of relational tables used in **SELECT** statement outer to the **inline view**



Correlated nested queries

In a **correlated nested query** an **inline view** may reference the names of relational tables used in **SELECT** statement outer to the **inline view**



Correlated nested queries

For example, consider a query: Find the chairs of departments that offer 12 credit point courses

Such query can be re-phrased as an equivalent query find the chairs of departments such that there exists at least one course worth 12 credit points offered by a department we are looking for

An inner query ... course worth 12 credit points offered by a department we are looking for is implemented as the following inline view

```
( SELECT *  
  FROM COURSE  
 WHERE credits = 12 AND offered_by = DEPARTMENT.name ) Q
```

Inline view

An outer query find the chairs of departments such that there exists at least one course found in the inner quer is implemented in the folowing way

```
SELECT chair  
FROM DEPARTMENT  
WHERE EXISTS Q
```

Query with an existential quantifier EXISTS that references and inline view

Correlated nested queries

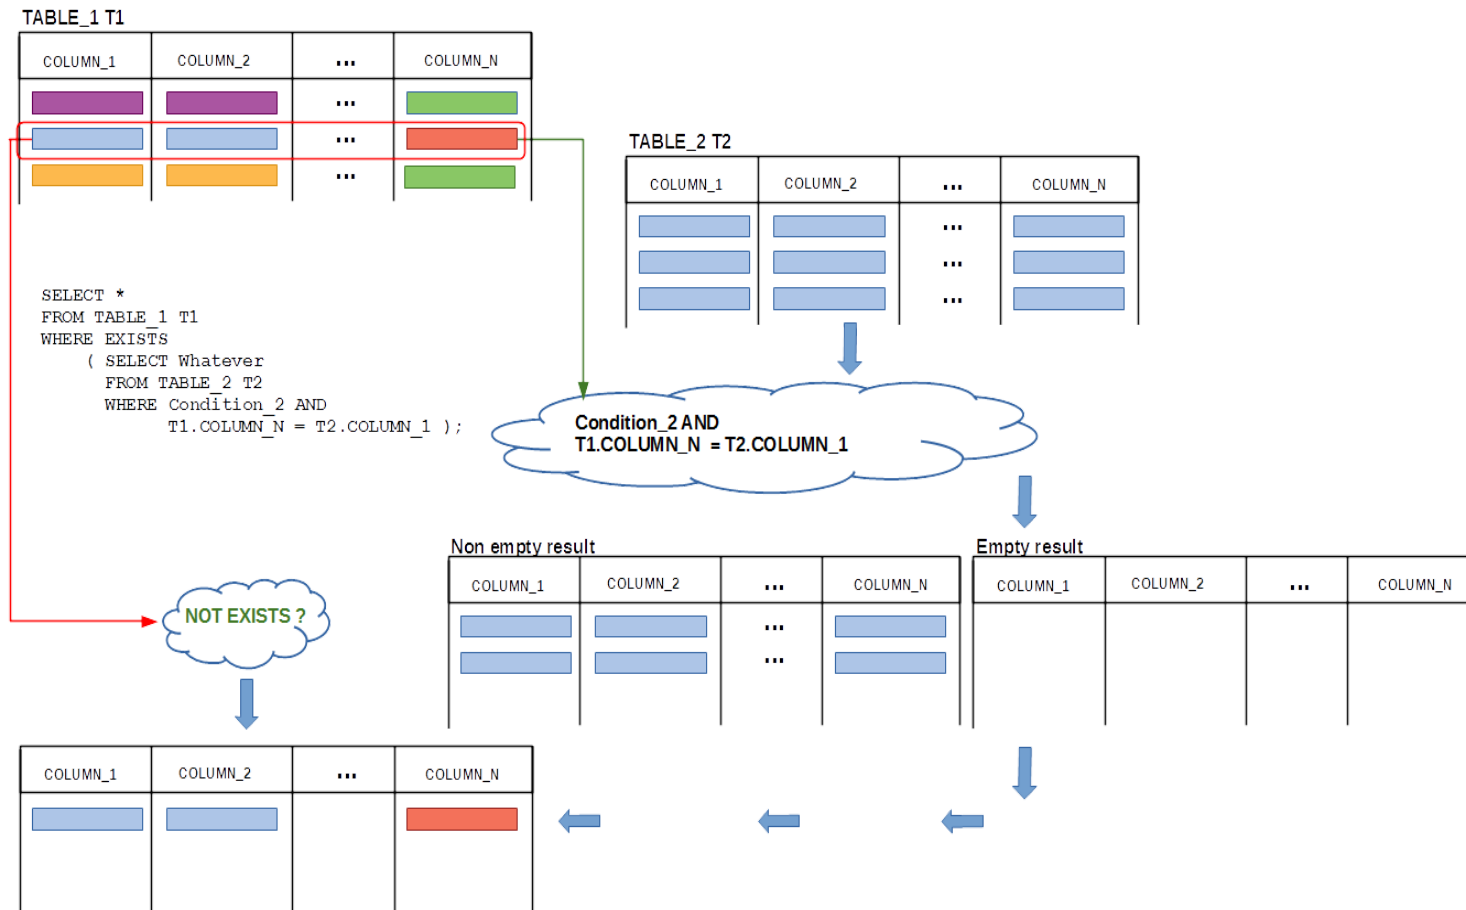
Finally, we replace a reference Q to an **inline view** with the **inline view** itself

Correlated nested query with an existential quantifier EXISTS

```
SELECT chair
FROM DEPARTMENT
WHERE EXISTS ( SELECT *
                FROM COURSE
                WHERE credits = 12 AND offered_by = DEPARTMENT.name );
```

Correlated nested queries

In a **correlated nested query** an **inline view** may reference the names of relational tables used in **SELECT** statement outer to the **inline view**



Correlated nested queries

Another example where **inline view** references a name of relational table used in **SELECT** statement outer to the **inline view** is the following

Find the chairs of all departments that offer no courses

It is equivalent to a query find the chairs of departments such that does not exist a course offered by a department we are looking for

An inner query finds ... a course offered by a department we are looking for

```
( SELECT *  
  FROM COURSE  
 WHERE offered_by = DEPARTMENT.name ) Q;
```

Inline view

An outer query finds all chairs of departments such that does not exist a course found by an inner query

```
SELECT chair  
FROM DEPARTMENT  
WHERE NOT EXISTS Q
```

Query with a negated existential quantifier EXISTS that references an inline view

Correlated nested queries

Finally, we replace a reference to an **inline view** Q with the **inline view** itself

Correlated nested query with a negated existential quantifier EXISTS

```
SELECT chair
FROM DEPARTMENT
WHERE NOT EXISTS( SELECT *
                  FROM COURSE
                  WHERE offered_by = DEPARTMENT.name );
```

SELECT statement (5)

Outline

Nested queries

Correlated nested queries

Queries with WITH clause

Relational views

Advanced DML statements

Queries with **WITH** clause

Consider a query: find the names of all departments together with the total number of courses offered by each department, include the departments that offer no courses

The query can be decomposed into the following two queries:

- find the names of departments and the numbers of courses offered by each department, and a query
- aggregate the results from the first query over the names of departments and count the total number of courses offered by each department

The first query can be implemented as a query definition **DEPT_COURSE** within **WITH** clause

```
WITH DEPT_COURSE AS
    ( SELECT name, cnum
      FROM DEPARTMENT LEFT OUTER JOIN COURSE
        ON DEPARTMENT.name = COURSE.offered_by ),
```

WITH clause with a query definition

Queries with WITH clause

The second query can be implemented a query definition `DC_COUNT` that references a query definition `DEPT_COURSE` within `WITH` clause

WITH clause with two query definitions

```
WITH DEPT_COURSE AS
    ( SELECT name, cnum
      FROM DEPARTMENT LEFT OUTER JOIN COURSE
        ON DEPARTMENT.name = COURSE.offered_by ),
  DC_COUNT AS
    ( SELECT name, COUNT(cnum) total_courses
      FROM DEPT_COURSE
      GROUP BY name )
```

Queries with WITH clause

The final query is implemented as `SELECT` statement appended to a query definition `DC_COUNT` within `WITH` clause

WITH clause with two query definitions and SELECT statement

```
WITH DEPT_COURSE AS
    ( SELECT name, cnum
      FROM DEPARTMENT LEFT OUTER JOIN COURSE
        ON DEPARTMENT.name = COURSE.offered_by ),
  DC_COUNT AS
    ( SELECT name, COUNT(cnum) total_courses
      FROM DEPT_COURSE
      GROUP BY name )

SELECT name
FROM DC_COUNT);
```


Queries with WITH clause

In another example **WITH** clause can be used to reduce the implementation complexity of the following query: **find the chairs of departments that offer both 6 and 12 credit point courses**

The query is decomposed into the following subqueries:

- Find the names of departments that offer 6 credit point courses
- Find the names of departments that offer 12 credit point courses
- Find the names of departments included in both results from the subqueries above
- Find the chairs of departments included in the previous subquery

The first subquery is implemented as the following query definition **COURSE6CR** within **WITH** clause

```
WITH COURSE6CR AS
    ( SELECT offered_by
      FROM COURSE
     WHERE credits = 6 ),
```

WITH clause with a query definition

Queries with **WITH** clause

A subquery: **find the chairs of departments that offer 12 credit point courses** is implemented as a query definition **COURSE12CR** appended to **WITH** clause

```
WITH COURSE6CR AS
    ( SELECT offered_by
      FROM COURSE
      WHERE credits = 6 ),
  COURSE12CR AS
    ( SELECT offered_by
      FROM COURSE
      WHERE credits = 12 );
```

WITH clause with two query definitions

Queries with WITH clause

A subquery: find the names of departments that offer both 6 and 12 credit point courses is implemented as a query definition
COURSE6_12CR appended to WITH clause

WITH clause with three query definitions

```
WITH COURSE6CR AS
    ( SELECT offered_by
      FROM COURSE
      WHERE credits = 6 ),
  COURSE12CR AS
    ( SELECT offered_by
      FROM COURSE
      WHERE credits = 12 ),
  COURSE6_12CR AS
    ( SELECT COURSE6CR.offered_by
      FROM COURSE6CR JOIN COURSE12CR
        ON COURSE6CR.offered_by = COURSE12CR.offered_by ),
```

Queries with WITH clause

A subquery: find the chairs of departments that offer both 6 and 12 credit point courses is implemented as a query definition CHAIR appended to WITH clause

WITH clause with four query definitions

```
WITH COURSE6CR AS
    ( SELECT offered_by
      FROM COURSE
      WHERE credits = 6 ),
  COURSE12CR AS
    ( SELECT offered_by
      FROM COURSE
      WHERE credits = 12 ),
  COURSE6_12CR AS
    ( SELECT COURSE6CR.offered_by
      FROM COURSE6CR JOIN COURSE12CR
        ON COURSE6CR.offered_by = COURSE12CR.offered_by ),
  CHAIR AS
    ( SELECT DEPARTMENT.chair
      FROM COURSE6_12CR JOIN DEPARTMENT
        ON COURSE6_12CR.offered_by = DEPARTMENT.name )
```

Queries with WITH clause

Finally a query: find the chairs of departments that offer both 6 and 12 credit point courses is implemented as `SELECT` statement appended to `WITH` clause

WITH clause with four query definitions and `SELECT` statement

```
WITH COURSE6CR AS
    ( SELECT offered_by
      FROM COURSE
      WHERE credits = 6 ),
  COURSE12CR AS
    ( SELECT offered_by
      FROM COURSE
      WHERE credits = 12 ),
  COURSE6_12CR AS
    ( SELECT COURSE6CR.offered_by
      FROM COURSE6CR JOIN COURSE12CR
        ON COURSE6CR.offered_by = COURSE12CR.offered_by ),
  CHAIR AS
    ( SELECT DEPARTMENT.chair
      FROM COURSE6_12CR JOIN DEPARTMENT
        ON COURSE6_12CR.offered_by = DEPARTMENT.name )

SELECT *
FROM CHAIR;
```

SELECT statement (5)

Outline

Nested queries

Correlated nested queries

Queries with WITH clause

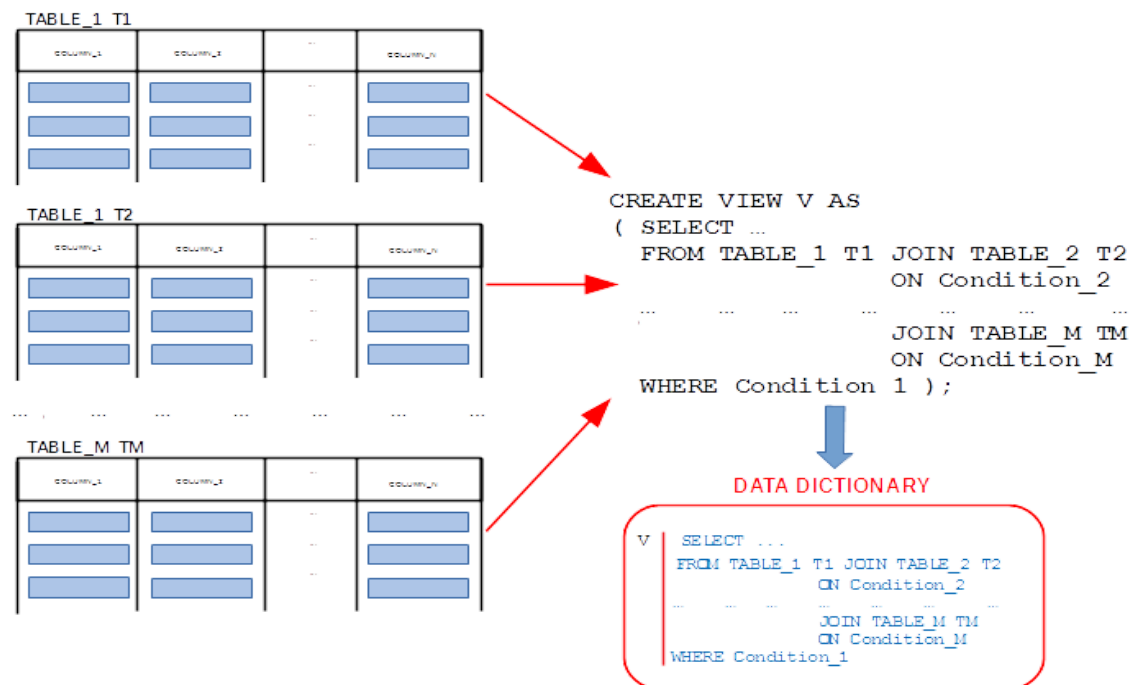
Relational views

Advanced DML statements

Relational views

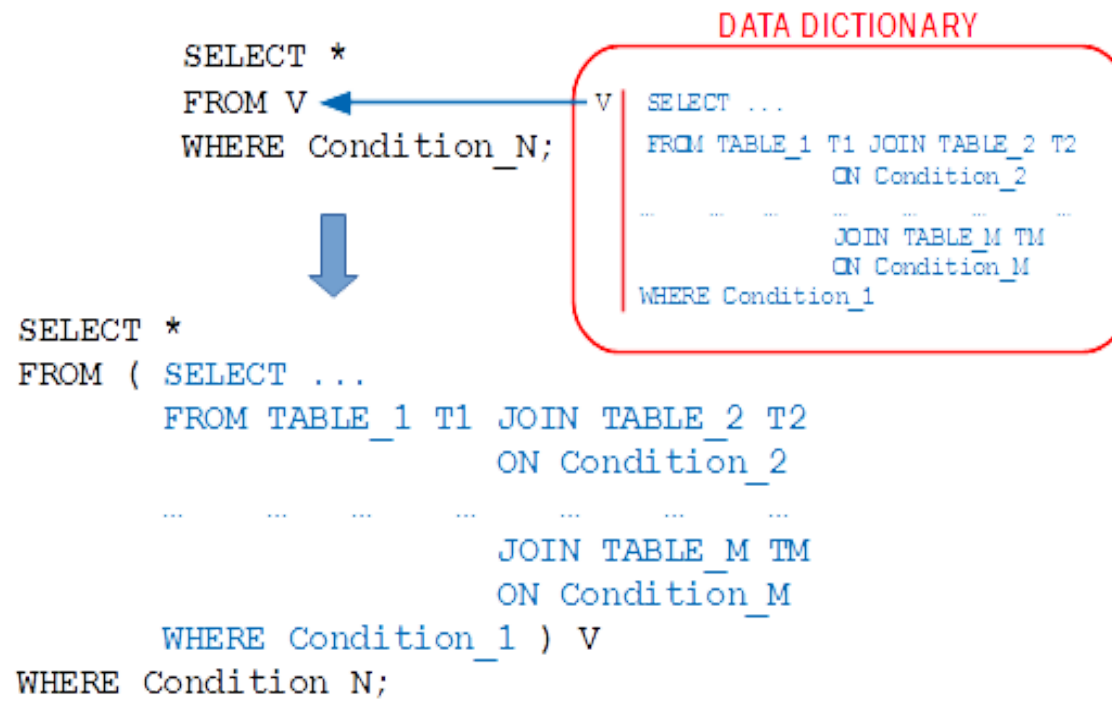
A **relational view** is a **virtual relational table** (derived relational table) that occupies no persistent storage and it is computed from very beginning each time it is used in **SELECT** statement

A **relational view** is stored by a database management system as a pair (name of a view, **SELECT** statement that defines the structure and contents of the view)



Relational views

Each time a name of **relational view** is used in **SELECT** statement its definition replaces the name of view and it becomes an **inline view**



Relational views

For example, create a relational view that contains the names of all departments together with the total number of courses offered by each department

CREATE VIEW statement

```
CREATE VIEW VDEPT( name, total_courses ) AS
( SELECT name, count(cnum)
  FROM DEPARTMENT LEFT OUTER JOIN COURSE
    ON DEPARTMENT.name = COURSE.offered_by
  GROUP BY name );
```

Relational views

Then a view `VDEPT` can be used to implement a query **find the names of departments that offer more than 1 course**

```
SELECT name
FROM VDEPT
WHERE total_courses > 1;
```

SELECT statement

The same query implemented with `GROUP BY` and `HAVING` clauses is the following

```
SELECT name, count(cnum)
FROM DEPARTMENT LEFT OUTER JOIN COURSE
                ON DEPARTMENT.name = COURSE.offered_by
GROUP BY name
HAVING count(cnum) > 1;
```

SELECT statement with LEFT OUTER JOIN operation

Relational views

A **relational view** can be used to reduce the complexity of **SELECT** statements

For example, a query **find the chairs of departments that offer both 6 and 12 credit point courses** is decomposed into the following queries

V6: Find the names of departments that offer 6 credit point courses

V12: Find the names of departments that offer 12 credit point courses

VNAME: Find the names of departments included in both **V6** and **V12**

VCHAIR: Find the chairs of departments included in **VNAME**

The views are implemented in the following way

V6: Find the names of departments that offer 6 credit point courses

```
CREATE VIEW V6( name ) AS
( SELECT offered_by
  FROM COURSE
 WHERE credits = 6 );
```

CREATE VIEW statement

Relational views

V12: Find the names of departments that offer 12 credit point courses

```
CREATE VIEW V12( name ) AS
( SELECT offered_by
  FROM COURSE
  WHERE credits = 12 );
```

CREATE VIEW statement

VNAME: Find the names of departments included in both V6 and V12

```
CREATE VIEW VNAME( name ) AS
( SELECT V6.name
  FROM V6 JOIN V12
    ON V6.name = V12.name );
```

CREATE VIEW statement

VCHAIR: Find the chairs of departments included in VNAME

```
CREATE VIEW VCHAIR( chair ) AS
( SELECT DEPARTMENT.chair
  FROM VNAME JOIN DEPARTMENT
    ON VNAME.name = DEPARTMENT.NAME );
```

CREATE VIEW statement

Relational views

The final query is

```
SELECT *  
FROM VCHAIR;
```

SELECT statement

SELECT statement (5)

Outline

Nested queries

Correlated nested queries

Queries with WITH clause

Relational views

Advanced DML statements

Advanced DML statements

Advanced Data Manipulation statements use subqueries implemented as `SELECT` statements inside DML statements

For example, delete all courses offered by a department chaired by Peter

DELETE statement with nested SELECT statement

```
DELETE FROM COURSE
WHERE offered_by = ( SELECT name
                     FROM DEPARTMENT
                     WHERE chair = 'Peter' );
```

Delete all departments that offer no courses

DELETE statement with negated existential quantifier EXISTS and nested SELECT statement

```
DELETE FROM DEPARTMENT
WHERE NOT EXISTS ( SELECT '1'
                  FROM COURSE
                  WHERE COURSE.offered_by = DEPARTMENT.name );
```

Advanced DML statements

Increase the total number of staff members by 5 in all departments that offer more than 20 courses

```
UPDATE DEPARTMENT
SET total_staff_number = total_staff_number + 5
WHERE name IN ( SELECT offered_by
                 FROM COURSE
                 GROUP BY offered_by
                 HAVING COUNT(cnum) > 20 );
```

UPDATE statement with nested SELECT statement

Add to table DEPARTMENT a column that contains the total number of courses offered by each department and insert the correct values into the column

```
ALTER TABLE DEPARTMENT ADD ( total_courses DECIMAL(2) );
```

ALTER TABLE statement that adds an attribute

```
UPDATE DEPARTMENT
SET total_courses = ( SELECT COUNT(title)
                     FROM COURSE
                     WHERE COURSE.offered_by = DEPARTMENT.name );
```

UPDATE statement with nested SELECT statement

Advanced DML statements

Create a table that contains the names of departments together with the total number of courses offered by each department and insert correct data into the table

```
CREATE TABLE DCNT AS
( SELECT name, COUNT(cnum) TOTC
  FROM DEPARTMENT LEFT OUTER JOIN COURSE
    ON DEPARTMENT.name = COURSE.offered_by
 GROUP BY name );
```

CREATE TABLE statement with nested SELECT statement

Note, that **NONE** of the consistency constraints except **NULL/NOT NULL** constraints are copied from the relational tables **DEPARTMENT** and **COURSE** into a relational table **DCNT**

Advanced DML statements

To preserve the consistency constraints we create a relational table
DCNT first ...

CREATE TABLE statement

```
CREATE TABLE DCNT(  
  name          VARCHAR(50)  NOT NULL,  
  total_courses DECIMAL(2)    NOT NULL,  
  CONSTRAINT DCNT_pkey PRIMARY KEY(name),  
  CONSTRAINT DCNT_fkey FOREIGN KEY (name) REFERENCES DEPARTMENT(name) );
```

... and we load data into the table next

INSERT statement with nested SELECT statement

```
INSERT INTO DCNT  
( SELECT name, COUNT(cnum)  
  FROM DEPARTMENT LEFT OUTER JOIN COURSE  
    ON DEPARTMENT.name = COURSE.offered_by  
 GROUP BY name );
```

References

T. Connolly, C. Begg, Database Systems, A Practical Approach to Design, Implementation, and Management, Chapter 6.3.5 Subqueries, Chapter 6.3.8 EXISTS and NOT EXISTS, Chapter 7.4.1 Creating a View Pearson Education Ltd, 2015

D. Darmawikarta, SQL for MySQL A Beginner's Tutorial, Chapter 7 Subqueries, Chapter 9 Views, Brainy Software Inc. First Edition: June 2014

[How to ... ? Cookbook, How to implement queries in SQL ? \(Part 2\) Recipe 6.3 How to implement nested queries ?](#)

[How to ... ? Cookbook, How to implement queries in SQL ? \(Part 2\) Recipe 6.4 Recipe 6.4 How to implement correlated nested queries ?](#)